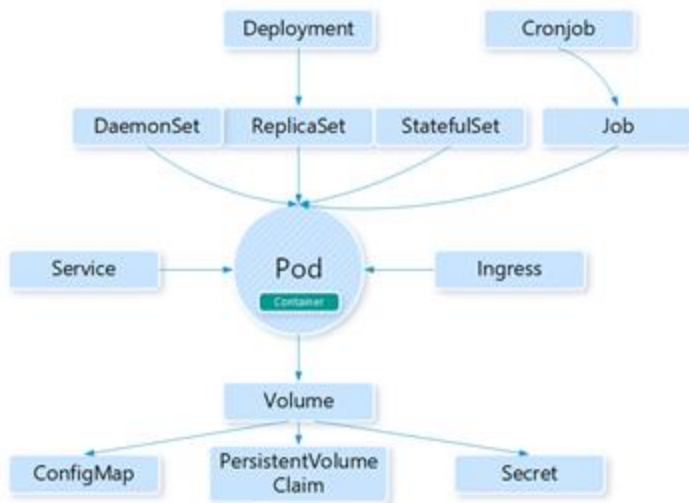


# Kubernetes Essential



- Assignment Review & Guides
- Kubernetes concepts
- Working with Pod
- Lab: Deployment Rolling Update and Rollback

- Everything within Kubernetes is as **an API Object**.
- Referenced within an object as the **apiVersion** and **kind**.



```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx:alpine
    name: nginx
    resources: {}
status: {}
```

- Objects represent the desired state of the object within the cluster.
- All objects **required**:
  - **apiVersion**: version of the Kubernetes API to create the object
  - **kind**: kind of object to create
  - **metadata**: data that helps uniquely identify the object, including a name string, UID, and optional namespace.
  - **spec**: desired state of the object

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx:alpine
    name: nginx
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

- Files or other representations of Kubernetes Objects are generally represented in YAML.
- Three basic data types:
  - List
  - Map
  - String, number, boolean, etc

```
# nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    app: nginx
    name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
```

# YAML and JSON

```
# nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  labels:
    app: nginx
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
```

```
# nginx-pod.json
{
  "kind": "Pod",
  "apiVersion": "v1",
  "metadata": {
    "name": "nginx",
    "labels": {
      "app": "nginx"
    }
  },
  "spec": {
    "containers": [
      {
        "name": "nginx",
        "image": "nginx",
      }
    ],
  },
}
```

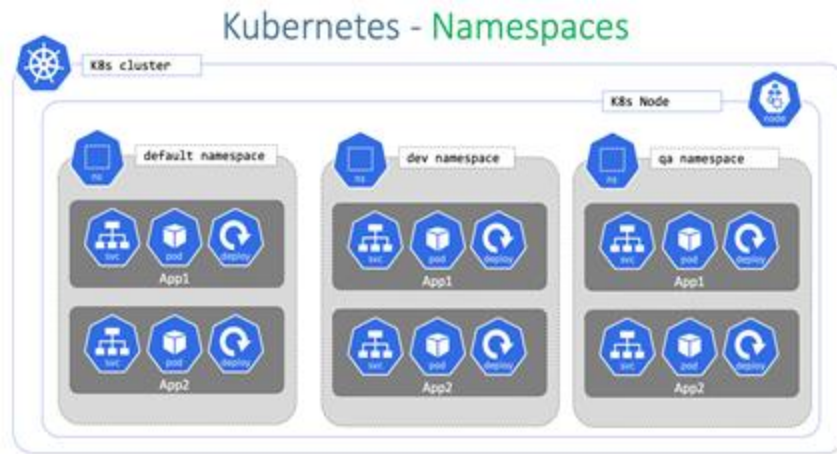
# CORE CONCEPTS

*namespaces* is used to isolate groups of resources within a single cluster:

- Names of resources need to be unique within a namespace, but not across namespaces.
- There are namespaced objects and cluster-wide objects.


```
[user@mate ~]$ kubectl create ns my-namespace  
namespace/my-namespace created
```

```
[user@mate ~]$ kubectl get ns  
NAME                STATUS   AGE  
default              Active   71m  
kube-node-lease      Active   71m  
kube-public          Active   71m  
kube-system          Active   71m  
local-path-storage   Active   71m
```





- Smallest unit of work of Kubernetes.
- Pods are one or MORE containers that share volumes, a network namespace, and are a part of a **single context**.



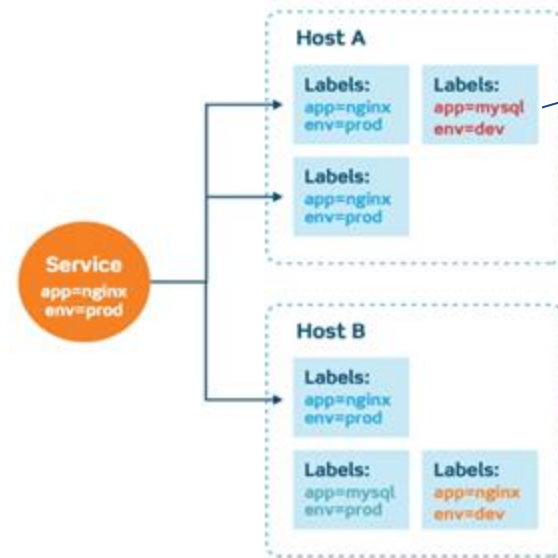
```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    app: demo2-01
  name: demo2-network
spec:
  containers:
    - image: redis:alpine
      name: redis
      resources: {}
    - image: busybox
      name: client
      command: ["nc"]
      args: ["-zv", "localhost", "6379"]
      resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

```
user@mate:~$ kubectl run new-nginx --image nginx --namespace my-ns
pod/new-nginx created
```

```
user@mate:~$ kubectl get pod --namespace my-ns
```

NAME	READY	STATUS	RESTARTS	AGE
new-nginx	1/1	Running	0	111s


- *Labels* are key-value pairs that are used to identify, describe and group together related sets of objects or resources.



```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    app: mysql
    env: dev
  name: demo2-network
spec:
  containers:
  - image: mysql
    name: mysql
    resources: {}
  - image: busybox
    name: client
    command: ["nc"]
    args: ["-zv", "localhost", "3306"]
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

*Selectors* use *labels* to filter or select objects, and are used throughout Kubernetes.

```
apiVersion: v1
kind: Pod
metadata:
  name: mlapp
  labels:
    app: mlapp
    env: prod
spec:
  nodeSelector:
    - gpu: nvidia
  containers:
    - name: nginx
      image: tensorflow/tensorflow
```

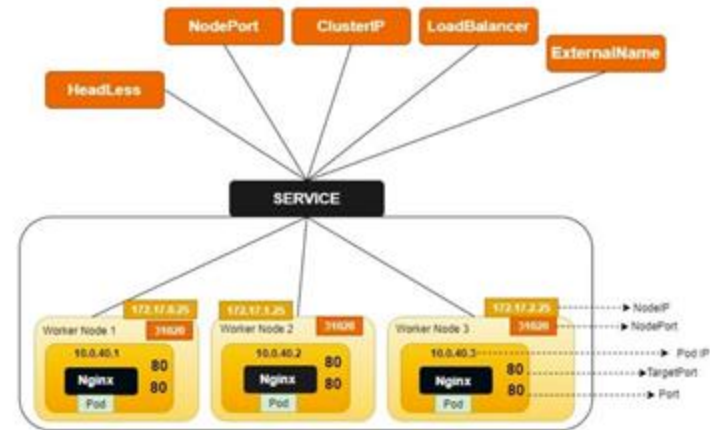


**Service** is a method for exposing a network application that is running as one or more Pods in the cluster:

- unique static cluster-wide IP
- static namespaced DNS name


There are 4 types:

- ClusterIP
- NodePort
- LoadBalancer
- ExternalName



**<service name>.<namespace>.svc.cluster.local**

*ClusterIP service* exposes the Service on a cluster-internal IP, uses labels and selectors to associate with pods. This Service only reachable from within the cluster.



```
# Create a ClusterIP service in my-ns namespace with port 8080 map to port 80
thainm5hls@xps:~$ kubectl create service clusterip --namespace my-ns new-cs --
tcp=8080:80
service/new-cs created

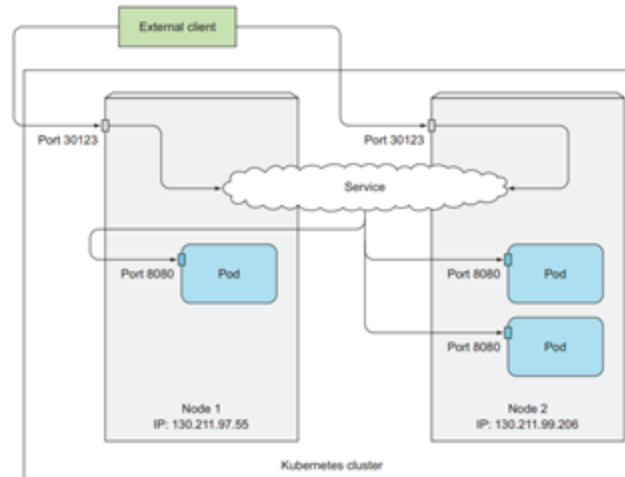
# Edit created service
thainm5hls@xps:~$ kubectl edit -n my-ns svc new-cs
```

```
# Get list of services in my-ns namespace
thainm5hls@xps:~$ kubectl get svc -n my-ns
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
new-cs	ClusterIP	10.96.163.21	<none>	8080/TCP	8m17s

```
# Get detailed information of service new-cs in namespace my-ns
thainm5hls@xps:~$ kubectl describe svc -n my-ns new-cs
```

- *NodePort service* makes Kubernetes reserve a port on all its nodes (the same port number) and forward incoming connections to the pods that are part of the service.
- Similar to a ClusterIP service, but a NodePort service can be accessed not only through the service's internal cluster IP, but also through any node's IP and the reserved node port.





```
# Create NodePort service nodeport-svc in namespace my-ns with service port 8080 map to backend port 80
thainm5hls@xps:~$ kubectl create service nodeport --namespace my-ns nodeport-svc --tcp=8080:80
service/nodeport-svc created
```

```
# List all services in namespace my-ns
thainm5hls@xps:~$ kubectl get svc -n my-ns
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
new-cs	ClusterIP	10.96.163.21	<none>	8080/TCP	23m
nodeport-svc	NodePort	10.96.197.227	<none>	8080:30175/TCP	9s

```
# Get detailed information of created NodePort service
thainm5hls@xps:~$ kubectl describe svc -n my-ns nodeport-svc
```

- Kubernetes clusters running on cloud providers usually support the automatic provision of a load balancer from the cloud infrastructure.
- The load balancer will have unique publicly accessible IP address and will redirect all connections to service.
- If Kubernetes is running in an environment that doesn't support LoadBalancer services, the load balancer will not be provisioned, but the service will still behave like a NodePort service.

# WORKLOADS

- Workloads within Kubernetes are higher level objects that manage Pods or other higher level objects.
- In **ALL CASES** a Pod Template is included, and acts the base tier of management.
- Types of workload:
  - ReplicaSet
  - Deployment
  - DaemonSet
  - StatefulSet
  - Job
  - Cronjob

- *Pod templates* are Pod specs with limited metadata.
- Workloads use *Pod templates* to make actual pods.

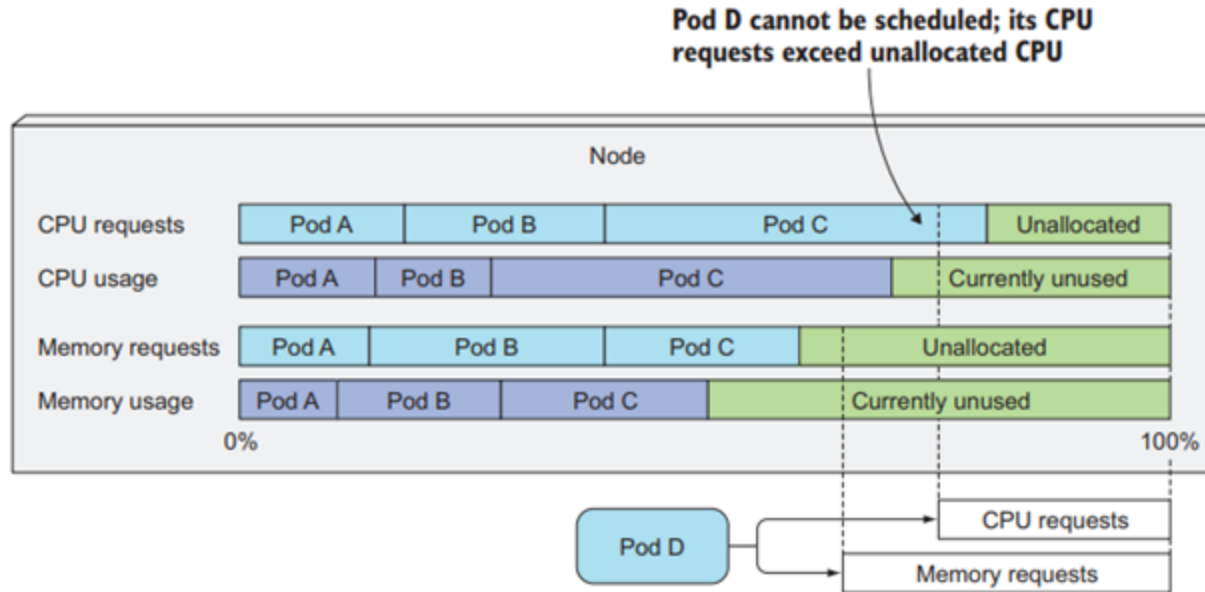
```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx:alpine
    name: nginx
    resources: {}
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: nginx
  name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
      - image: nginx:alpine
        name: nginx
        resources: {}
```

- When create a pod, the amount of CPU and memory that a container needs- *requests*, and a hard limit on what it may consume- *limits*, can be specified.
  - CPU unit: core
  - Memory unit: byte

The container request at least 200 milicore of CPU ( $\frac{1}{5}$  of a single core) and 100Mi (0.1M) of memory.

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    run: nginx
  name: nginx
spec:
  containers:
  - image: nginx
    name: nginx
    resources:
      requests:
        cpu: 200m
        memory: 100Mi
```



**Figure 14.1** The Scheduler only cares about requests, not actual usage.

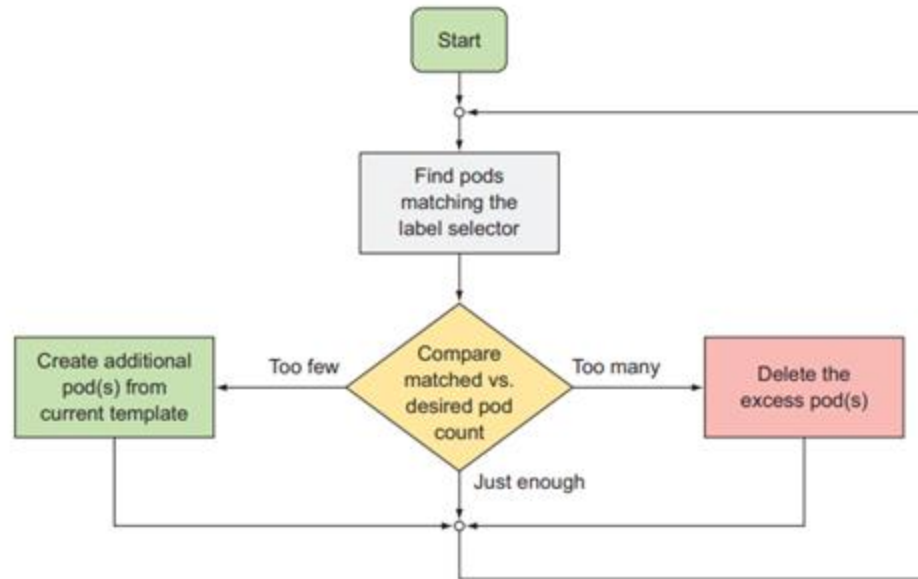
Init containers run before app container are started.

Init containers are exactly like regular containers, except:

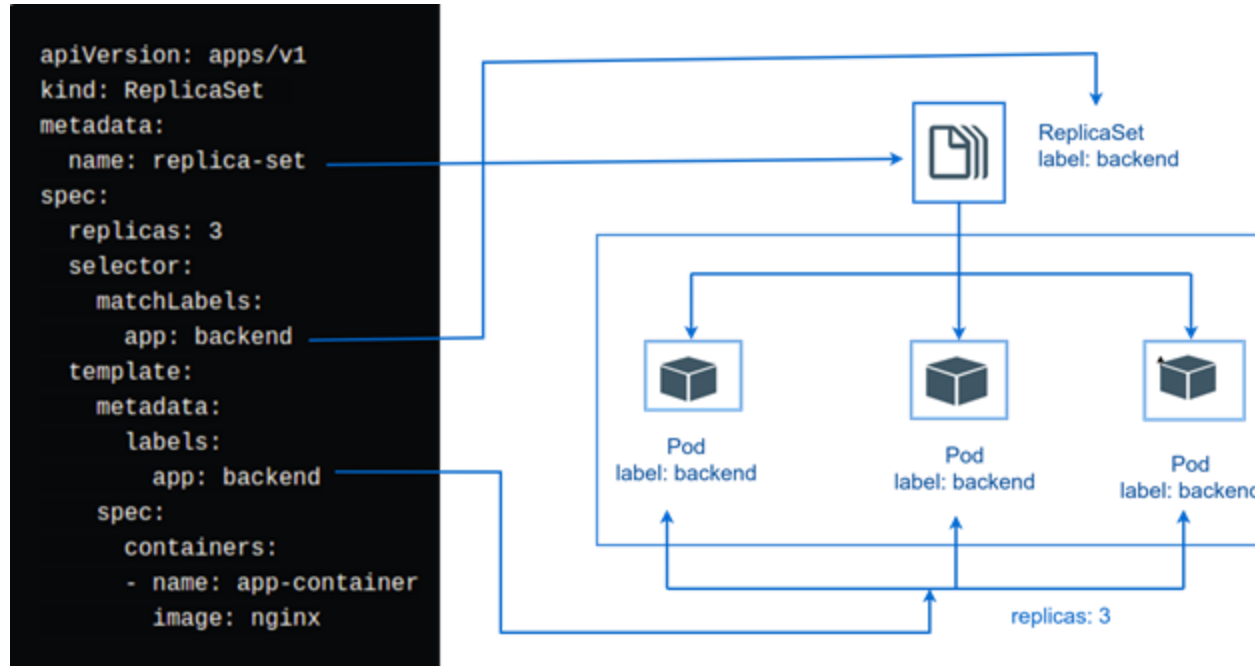
- Init containers always run to completion.
- Each init container must complete successfully before the next one starts.



- *ReplicaSet* ensures pods are always kept running.
- If the pod disappears for any reason, the *ReplicaSet* notices the missing pod and creates a replacement pod.



# ReplicaSet



- *Deployment* provides a declarative way to describe the desired state of your application. Kubernetes then takes care of creating and managing the necessary *ReplicaSets* and Pods to ensure that the desired state is achieved.
- Each update creates a *pod-template-hash* label by hashing *podTemplate* and then assign to both the *ReplicaSet* and subsequent *Pods*.



- *.spec.strategy*: specifies the strategy used to replace old Pods by new ones. Available options:
  - *Recreate*: All existing Pods are killed before new ones are created.
  - *RollingUpdate* (default): Updates Pods in a rolling update fashion.
- *.spec.revisionHistoryLimit*: specifies the number of old ReplicaSets to retain to allow rollback. By default, 10 old *ReplicaSets* will be kept.

```
thainm5hls@xps:~/$ kubectl create deployment nginx-deployment --image nginx:1.16 --replicas 3
deployment.apps/nginx-deployment created
```

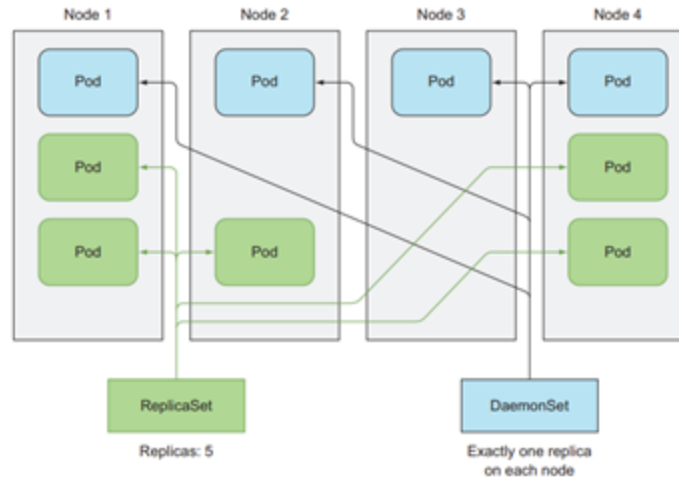
```
thainm5hls@xps:~/$ kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-deployment-66fff56987-gs9bs	1/1	Running	0	30s
nginx-deployment-66fff56987-nsnms	1/1	Running	0	30s
nginx-deployment-66fff56987-tgg9p	1/1	Running	0	30s

```
thainm5hls@xps:~/$ kubectl get deployments.apps
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-deployment	3/3	3	3	118s

- *DaemonSet* ensures there is exactly one instance of a pod running on each node, even with a new added node, and is ideal for cluster wide services such as log forwarding, or health monitoring.
- *DaemonSet* **skips** default scheduling mechanisms.



- **Job** ensures one or more pods are executed and successfully terminate.
- If :
  - The process running inside the Pod's container finish successfully, the pod is considered completed.
  - The process fails and returns an error exit code, the Job can be configured to either restart the container or not.
  - A node failure, the pods on that node that are managed by a Job will be rescheduled to other nodes.
- Pods are NOT cleaned up until the job itself is deleted.

- **Job** pods can't use the default policy, because they're not meant to run indefinitely. Therefore, you need to explicitly set the restart policy to either *OnFailure* or *Never*.
- Some helpful attributes:
  - *.spec.backoffLimit*: The number of failures before the job itself is considered failed.
  - *.spec.parallelism*: How many instances of the pod can be run concurrently.
  - *.spec.completion*: How many instances of the pod can be run concurrently.



- An extension of the Job Controller, it provides a method of executing jobs on a cron-like schedule.

